

Continuous Integration Timeouts, Automated Program Repair, and Crash Triage in Game Development: A PRISMA-Compliant Systematic Review and Meta-Analysis

Authors: Systematic Meta-Analysis Research Consortium | Protocol: PRISMA 2020 | Date: August 2026

> [SYNTHETIC PIPELINE DEMONSTRATION] This manuscript, including all data, citations, and statistics, is entirely synthetically generated by an AI as a pipeline scaffold. It does not represent real scientific findings.

Authors: [SYNTHETIC PIPELINE SCAFFOLD]

Affiliation: [SYNTHETIC DATA PROTOTYPE]

Protocol Reference: PRISMA-P 2020 Registration Protocol

Date: August 2026

Abstract

Background: Continuous Integration (CI) in video game engineering operates under stringent constraints fundamentally distinct from traditional enterprise software. Game CI pipelines frequently suffer from catastrophic build timeouts and un-triaged test crashes caused by three interrelated factors: (1) gigabyte-scale binary asset cooking and shader compilation bottlenecks, (2) physics and asynchronous frame-loop non-determinism, and (3) complex multi-language engine architectures (e.g., C++ core interacting with C#/Lua gameplay layers).

Objectives: To systematically identify, categorize, and synthesize empirical literature regarding CI timeout mitigation, crash triage, and automated program repair (APR) in interactive software and game engineering, evaluating their efficacy and identifying critical benchmark gaps.

Data Sources: Primary ingestion across CrossRef (ACM Digital Library records) and arXiv APIs yielded 482 raw records.

Eligibility Criteria: Studies were evaluated via a strict 5-stage PICOC protocol targeting automated program repair, crash reproduction, non-deterministic test mitigation, and CI scheduling in game engines, interactive software, or distributed systems.

Synthesis Methods: A random-effects meta-analysis (DerSimonian-Laird model with inverse-variance fixed-effects comparison) was conducted on standardized effect sizes (Hedges' g) for quantitative runtime reductions and triage speedups. Statistical heterogeneity was assessed via Cochran's Q and the I^2 index.

Results: Of 482 identified records, 50 passed Title/Abstract screening, and 28 met all full-text criteria. Across quantitative timing benchmarks ($k = 9$), automated interventions demonstrated a statistically significant pooled effect size of **Hedges' $g = 4.404$** (95% CI: [3.623, 5.185], $Z = 11.05$, $p < 0.001$). However, extreme statistical heterogeneity (**$I^2 = 89.1\%$** , $Q = 73.57$, $\tau^2 = 1.242$) was observed. Furthermore, a substantial **+50.0% domain disparity gap** exists: 71.4% of studies (20/28) evaluate exclusively on enterprise benchmarks (Defects4J), while only 21.4% (6/28) address game engines or interactive software.

Discussion & Conclusion: The observed heterogeneity ($I^2 = 89.1\%$) and domain disparity demonstrate that existing APR and CI tools are evaluated in sanitized, deterministic environments that fail to reflect the asset-heavy, non-deterministic reality of game development. This paper mathematically justifies the urgent deployment of **OSSGameBench** as a standardized, multi-language benchmark suite tailored to game engine CI pipelines.

Keywords: Continuous Integration, CI Timeouts, Automated Program Repair, Crash Triage, Game Engine Architecture, Flaky Tests, PRISMA 2020, OSSGameBench.


1. Introduction

Continuous Integration (CI) and Continuous Deployment (CD) pipelines have become the backbone of modern software quality assurance [Jin & Servant, 2023]. By automatically building, integrating, and testing source code modifications, CI systems allow engineering teams to detect regressions rapidly. However, within video game development and interactive 3D application engineering, CI pipelines encounter unique, systemic challenges that cause severe developer friction and high operational costs [Zhang et al., 2025].

1.1 The Triad of Game Development CI Failures

Game engineering CI pipelines fail or time out primarily due to three architectural dynamics:

- 1. **Massive Binary Assets & Cooking Pipelines:** Unlike traditional enterprise software consisting primarily of text source files, video games depend on gigabytes of binary assets, including high-resolution textures, 3D meshes, skeleton rigs, and shader graphs. Asset cooking—the transformation of source assets into platform-specific optimized runtime formats—is computationally expensive and I/O intensive. A full asset bake frequently exceeds standard CI runner timeout limits (typically 30 to 60 minutes), causing pipeline aborts unrelated to source-level code defects.
- 2. **Physics & Frame-Loop Non-Determinism:** Video games rely on continuous game loops driven by variable frame rates, asynchronous multithreaded rendering, and physics engines (e.g., PhysX, Box2D, Havok). Small floating-point precision differences and thread scheduling jitter produce non-deterministic test execution, resulting in 'flaky' test failures that masquerade as regression bugs [Yi et al., 2019; Lou et al., 2024].
- 3. **Cross-Language Stack Traces & Crash Reproduction:** Production game engines typically combine a high-performance native core (C/C++) with high-level scripting environments (C#, Lua, Python, Visual Scripting). When crashes occur in CI test runs, stack traces frequently cross the interop boundary, creating corrupt or opaque crash dumps that defeat conventional fault localization tools [Zhao et al., 2022].



1.2 The Empirical Benchmark Gap

Despite the high economic stakes of game development, automated program repair (APR) and CI optimization techniques have been overwhelmingly evaluated on standard enterprise datasets such as Defects4J [Wei et al., 2026], ManyBugs, and QuixBugs. These benchmarks consist of small, isolated Java or C algorithms with fast, deterministic unit test suites, completely omitting binary assets, graphics pipelines, and game loop timing dependencies.

1.3 Research Questions

- To systematically investigate this problem, this PRISMA-compliant meta-analysis addresses the following four Research Questions (RQs):
- **RQ1 (Intervention Scope):** What automated techniques (APR, crash triage, test selection, timeout mitigation) have been proposed to resolve CI failures and crashes?
 - **RQ2 (Empirical Efficacy):** What is the synthesized standardized effect size of automated interventions on CI execution time and crash triage latency?

- **RQ3 (Heterogeneity & Variance):** What degree of statistical heterogeneity exists across existing empirical studies, and what factors drive between-study variance?
 - **RQ4 (Domain Generalizability):** To what extent do existing benchmark evaluation environments represent game engine architectures, and how does this justify the need for dedicated suites such as OSSGameBench?
-

2. Methodology

This review was conducted in accordance with the **PRISMA 2020 (Preferred Reporting Items for Systematic Reviews and Meta-Analyses)** statement [Page et al., 2021].

2.1 Search Strategy & Data Ingestion

The literature search was executed using automated ingestion scripts querying primary academic APIs:

- **ACM Digital Library** via CrossRef REST API (filtering for Publisher Member ID 320 to capture flagship venues including ICSE, FSE, ASE, ISSTA, and TOG).
- **arXiv API** (covering computer science subcategories cs.SE, cs.PL, cs.GR).

The formalized Boolean search string was defined as:

Q = ('CI timeout' OR 'continuous integration') AND ('crash triage' OR 'automated repair' OR 'OSSGameBench') AND 'game'

A total of **482 distinct bibliographic records** were ingested. All Digital Object Identifiers (DOIs) were stripped of URL prefixes, sanitized to lowercase strings, and subjected to strict deduplication.

2.2 Screening Protocol (PICOC Framework)

Records underwent multi-stage screening based on the formal **PICOC (Population, Intervention, Comparison, Outcomes, Context)** specification:

- **Population (P):** Continuous integration pipelines, game software architectures, crash triage systems, or non-deterministic test suites.
- **Intervention (I):** Automated program repair (APR), automated crash localization/reproduction, flaky test mitigation, or predictive CI test/build selection.
- **Comparison (C):** Empirical baseline comparisons (e.g., standard CI runs, manual debugging, unguided LLMs, or random test generation).
- **Outcomes (O):** Actionable metrics including CI timeout reduction (%), triage precision/recall (%), patch generation success rate (%), or time-to-fix speedup.
- **Context (C):** Computational software engineering, game development, or interactive/distributed systems.

Fail-fast logic was enforced: any record failing a mandatory PICOC boolean rule was immediately classified as EXCLUDE.

2.3 Full-Text Retrieval & Eligibility

Candidate records passing Title/Abstract screening were subjected to full-text retrieval using direct open-access resolvers (Unpaywall API, OpenAlex API, and Semantic Scholar open access endpoints). Extracted texts and comprehensive bibliographic ledgers were screened for domain relevance across three dimensions:

1. *Game Engine Architecture:* Explicit focus on 2D/3D engines (Unreal, Unity, Godot, custom engines), game loops, and rendering subsystems.
2. *Massive Binary Assets:* Modeling binary asset cooking, Git LFS asset pipelines, or asset-induced CI timeout exhaustion.

3. *Non-Deterministic Testing & Crash Triage*: Modeling frame-rate flakiness, floating-point error cancellation, and automated triage.

2.4 Statistical Meta-Analysis Model

To synthesize quantitative empirical outcomes across studies reporting execution time savings and triage speedups, we computed standardized effect sizes using **Hedges' g** (Cohen's d with small-sample bias correction J):

$$d = (X_{\text{intervention}} - X_{\text{control}}) / S_{\text{pooled}}$$
$$J = 1 - (3 / (4 * (n1 + n2) - 9))$$
$$g = J * d$$
$$\text{Var}(g) = ((n1 + n2) / (n1 * n2)) + (g^2 / (2 * (n1 + n2)))$$

Meta-analytic pooling was computed using the **DerSimonian-Laird Random-Effects Model**:

$$w_i^* = 1 / (\text{Var}(g_i) + \tau^2)$$
$$g_{\text{pooled}} = \text{sum}(w_i * g_i) / \text{sum}(w_i^*)$$

Statistical heterogeneity was quantified using Cochran's Q and Higgins' I²:

$$Q = \text{sum}(w_i * (g_i - g_{\text{fixed}})^2)$$
$$I^2 = \max(0, (Q - (k-1)) / Q) * 100\%$$

3. Results

3.1 PRISMA Flow Breakdown

The identification and screening progression is summarized in the following table:

Phase	Description	Records	% of Screened
Identification	Records harvested from primary academic databases (CrossRef/ACM, arXiv)	482	100.0%
Deduplication	Unique records post-lowercase DOI normalization	482	100.0%
Title/Abstract Screening	Records screened against strict PICOC rules	482	100.0%
	<i>Excluded: Outside target population (P)</i>	246	51.0%
	<i>Excluded: No automated triage/repair intervention (I)</i>	91	18.9%
	<i>Excluded: Lacks empirical comparison baseline (C)</i>	81	16.8%
	<i>Excluded: Lacks actionable efficiency/accuracy metrics (O)</i>	14	2.9%
	<i>Candidate records promoted to Full-Text Eligibility</i>	50	10.4%
Full-Text Retrieval	Full-text PDFs downloaded & stored locally (`full_text_pdfs/`)	6	12.0%
	Extended bibliographic records & full abstracts evaluated	44	88.0%
Eligibility Assessment	Full texts evaluated for game engine, asset, & crash triage scope	50	100.0%
	<i>Excluded: Lacks explicit game engine/asset/crash relevance</i>	22	44.0%
Final Synthesis	Total approved studies synthesized in Taxonomy Matrix	28	5.8%

3.2 Taxonomy & Dimensional Breakdown (N = 28)

The 28 approved studies were classified across intervention categories and evaluation targets:

Intervention Classification:	
* Automated Program Repair (APR)	: 24 studies (85.7%)
* Crash Localization & Triage	: 13 studies (46.4%)
* Flaky Test & Non-Determinism Mitigation	: 6 studies (21.4%)
* CI Scheduler Optimization & Selection	: 3 studies (10.7%)
Evaluation Target Classification:	
* Defects4J / Standard SE Benchmarks	: 20 studies (71.4%)
* AAA Industrial & Open-Source Repos	: 12 studies (42.9%)
* Game Engines & Interactive Software	: 6 studies (21.4%)

3.3 Quantitative Statistical Meta-Analysis (k = 9)

Nine studies provided rigorous, extractable quantitative data on execution acceleration, timeout reduction, or turnaround speedup:

Study ID	Reference	Metric Focus	Hedges' g	95% Confidence Interval	Weight (%)
STUDY_25	Jin & Servant [2023]	CI Build Time / Timeout Reduction (%)	3.91	[3.30, 4.52]	11.85%
STUDY_02	Wei et al. [2026]	Patch Validity & Generation Speedup	4.00	[3.28, 4.71]	11.54%
STUDY_03	Zirak & Hemmati [2024]	Exact Match Fix Efficacy Gain (%)	2.96	[2.33, 3.59]	11.78%
STUDY_04	Zhang et al. [2025]	Equivalence Class Validation Speedup	4.03	[3.22, 4.85]	11.21%
STUDY_09	Zhao et al. [2022]	Crash Trace Reproduction Speedup	5.17	[4.35, 5.98]	11.21%
STUDY_19	Yi et al. [2019]	Numerical Instability Repair Speedup	6.91	[5.56, 8.27]	8.84%
STUDY_21	Kong et al. [2025]	Conversational Round Latency Reduction	2.59	[2.01, 3.17]	11.96%
STUDY_23	Ma et al. [2026]	Verification Time Reduction	5.84	[4.74, 6.94]	10.02%
STUDY_28	Liu et al. [2023]	Search Space & Timeout Reduction (%)	4.88	[4.15, 5.61]	11.59%
Pooled	Overall Synthesis	Combined Timing & Triage Efficiency	4.40	[3.62, 5.18]	100.0%

3.4 Forest Plot Synthesis

The random-effects meta-analysis demonstrated a highly significant overall effect size favoring automated interventions:
Pooled g = 4.404, 95% CI = [3.623, 5.185], Z = 11.05, p < 0.001

The high-resolution forest plot visualizing individual study effect sizes and the synthesized DerSimonian-Laird summary diamond is embedded below:




Figure 1: PRISMA Meta-Analysis Forest Plot: Standardized Effect Sizes (Hedges' g with 95% CI) and Random-Effects Synthesis.

4. Discussion

4.1 Statistical Heterogeneity and the Need for Standardized Environments

A central finding of this meta-analysis is the substantial statistical heterogeneity across the included literature:
 $I^2 = 89.1\%$, Cochran's $Q = 73.57$ (df = 8, $p < 0.0001$), $\tau^2 = 1.242$

In classical epidemiological meta-analysis, an $I^2 > 75\%$ indicates substantial heterogeneity. In software engineering, this degree of heterogeneity stems directly from **unstandardized experimental evaluation platforms**:

- Divergent Baseline Definitions:** Studies such as *Yi et al. [2019]* evaluate floating-point APR against unguided random perturbation ($g = 6.91$), whereas *Kong et al. [2025]* evaluate conversational LLM repair against advanced multi-turn prompting baselines ($g = 2.59$).
- Artificial Test Harness Isolation:** Tools are evaluated in synthetic test environments where compilation times and test execution are sub-second, masking the multi-minute overhead of industrial build runners.
- Absence of Unified Metrics:** While CI optimization studies [Jin & Servant, 2023] report wall-clock execution savings, APR studies report patch candidate counts, hindering direct cross-paradigm comparisons.

4.2 The +50.0% Domain Disparity Gap

Our analysis reveals an acute empirical discrepancy between software engineering research benchmarks and game development operational reality:

Standard SE Benchmark Representation (Defects4J) = $20 / 28 = 71.4\%$

Game Engines & Interactive Software Representation = $6 / 28 = 21.4\%$

Disparity Gap (Enterprise Bias) = 71.4% - 21.4% = +50.0%



This +50.0% gap highlights that current software repair tools are over-optimized for small, deterministic Java programs. When deployed in commercial game studios (e.g., Unreal Engine or Unity projects), these tools fail because they cannot handle:

- **Multi-gigabyte Git LFS repositories** with texture/mesh dependencies.
- **Shader compilation matrices** (DirectX, Vulkan, Metal) that bottleneck CI runners.
- **Frame-rate-dependent race conditions** and physics simulation variance.

4.3 Theoretical Formulation & Mathematical Justification for OSSGameBench

The combination of **extreme heterogeneity** ($I^2 = 89.1\%$) and the **+50.0% enterprise disparity** provides the mathematical justification for **OSSGameBench**:

Benchmarking Need Index (B) = $I^2 \text{ Disparity Gap} = 0.891 \cdot 0.500 = 0.4455$

OSSGameBench directly resolves these deficiencies by establishing:

1. **Real-World Open-Source Game Engine Targets:** Incorporating mature engines (Godot Engine, OpenMW, Veloren, Panda3D) with production-grade build scripts.
2. **Reproducible Non-Deterministic Test Suites:** Providing physics and rendering benchmark harnesses with statistical flake-filtering oracles.
3. **Standardized CI Timeout Scenarios:** Benchmarking asset cooking pipelines and test selection algorithms under realistic runner resource constraints.

4.4 Threats to Validity

- **Internal Validity:** Variability in reported quantitative parameters across papers required estimating pooled standard deviations from empirical distributions. Sensitivity analyses confirmed that minor variance perturbations do not alter the statistical significance ($p < 0.001$).
- **External Validity:** Only English-language academic papers were ingested, potentially excluding internal proprietary grey literature from closed AAA game studios.
- **Construct Validity:** The PICOC protocol used keyword matching alongside fail-fast semantic filtering; some domain studies that avoided explicit 'continuous integration' keywords may have been excluded.

5. Practical Implications & Conclusion

5.1 Recommendations for Game AI Engineers and Researchers

1. **Adopt Combined Build & Test Selection in CI:** Implement predictive selection algorithms (e.g., *HybridCISave*) to prioritize impacted compilation units and reduce asset baking overhead by up to 68.3%.
2. **Employ Contrastive Test Oracles for Flaky Game Physics:** Use contrastive passing/failing telemetry pairs (*ContrastRepair*) to isolate rendering race conditions from legitimate logic defects.

3. **Standardize on OSSGameBench:** Future research in automated repair and CI optimization for interactive software must evaluate against OSSGameBench to ensure generalizability to multi-language, asset-heavy game development.

5.2 Conclusion

This PRISMA systematic review and meta-analysis synthesized 28 peer-reviewed studies investigating CI timeouts, crash triage, and automated program repair. Automated interventions demonstrate a powerful pooled effect size ($g = 4.404$, $p < 0.001$) in accelerating bug fixes and reducing build overhead. However, the substantial heterogeneity ($I^2 = 89.1\%$) and +50.0% enterprise benchmark bias confirm that current research remains disconnected from game engineering realities. The adoption of OSSGameBench provides a vital path forward for rigorous, reproducible software engineering in game development.

References

1. Bonakdarpour, B., & Kulkarni, S. S. (2012). Automated model repair for distributed programs. *ACM Transactions on Computer Systems*, DOI: 10.1145/2261417.2261437.
2. Gao, X., Wang, B., Duck, G. J., Ji, R., Xiong, Y., & Roychoudhury, A. (2021). Beyond Tests: Program Repair with Formal Specifications. *ACM Transactions on Software Engineering and Methodology*, DOI: 10.1145/3418461.
3. Hossain, S. B., Jiang, N., Zhou, Q., Li, X., Chiang, W. H., Lyu, Y., Nguyen, H., & Tripp, O. (2024). A Deep Dive into Large Language Models for Automated Bug Localization and Repair. *ACM FSE*, DOI: 10.1145/3660773.
4. Huang, K., Xu, Z., Yang, S., Sun, H., Li, X., Yan, Z., & Zhang, Y. (2024). Evolving Paradigms in Automated Program Repair: Taxonomy, Challenges, and Opportunities. *ACM Computing Surveys*, DOI: 10.1145/3696450.
5. Intarasirisawat, J., Ang, C. S., Efstratiou, C., Dickens, L., Sriburapar, N., Sharma, D., & Asawathaweboon, B. (2020). An Automated Mobile Game-based Screening Tool for Clinical State Classification. *ACM IMWUT*, DOI: 10.1145/3411837.
6. Jin, X., & Servant, F. (2023). HybridCISave: A Combined Build and Test Selection Approach in Continuous Integration. *ACM ISSTA*, DOI: 10.1145/3576038.
7. Kong, J., Xie, X., Cheng, M., Liu, S., Du, X., & Guo, Q. (2025). ContrastRepair: Enhancing Conversation-Based Automated Program Repair via Contrastive Test Case Pairs. *ACM TOSEM*, DOI: 10.1145/3719345.
8. Li, F., Jiang, J., Sun, J., & Zhang, H. (2025). Hybrid Automated Program Repair by Combining Large Language Models and Program Analysis. *ACM TOSEM*, DOI: 10.1145/3715004.
9. Liu, K., Zhang, J., Li, L., Koyuncu, A., Kim, D., Ge, C., Liu, Z., Klein, J., & Bissyandé, T. F. (2023). Reliable Fix Patterns Inferred from Static Checkers for Automated Program Repair. *ACM TOSEM*, DOI: 10.1145/3579637.
10. Lou, Y., Yang, J., Benton, S., Hao, D., Tan, L., Chen, Z., Zhang, L., & Zhang, L. (2024). When Automated Program Repair Meets Regression Testing: An Extensive Study on Two Million Patches. *ACM TOSEM*, DOI: 10.1145/3672450.
11. Ma, C., Shi, J., Han, R., Liu, Y., Li, F., Niu, Y., & Lo, D. (2026). Automated Repair of TEE Partitioning Issues via DSL-Guided and LLM-Assisted Patching. *ACM ICSE*, DOI: 10.1145/3808207.
12. Martinez, M., Martínez-Fernández, S., & Franch, X. (2025). The Sustainability Face of Automated Program Repair Tools. *ACM TOSEM*, DOI: 10.1145/3744900.
13. Page, M. J., et al. (2021). The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. *BMJ*, 372:n71.
14. Qi, Y., Mao, X., Lei, Y., & Wang, C. (2013). Using automated program repair for evaluating the effectiveness of fault localization techniques. *ACM ISSTA*, DOI: 10.1145/2483760.2483785.

15. **Schulte, E., DiLorenzo, J., Weimer, W., & Forrest, S.** (2013). Automated repair of binary and assembly programs for cooperating embedded devices. *ACM ASPLOS*, DOI: 10.1145/2499368.2451151.
16. **Vallecillos-Ruiz, F., Grishina, A., Hort, M., & Moonen, L.** (2026). Assessing the Latent Automated Program Repair Capabilities of Large Language Models Using Round-Trip Translation. *ACM TOSEM*, DOI: 10.1145/3771922.
17. **Wei, Y., Xia, C. S., & Zhang, L.** (2026). Copiloting the Copilots for Automated Program Repair. *ACM ICSE*, DOI: 10.1145/3788082.
18. **Xiang, J., Xu, X., Kong, F., Wu, M., Zhan, Z., Zhang, H., & Zhang, Y.** (2026). Practical LLM-Based Function-Level Automated Program Repair: How Far Are We? *ACM ICSE*, DOI: 10.1145/3812804.
19. **Yang, D., Lei, Y., Mao, X., Qi, Y., & Yi, X.** (2023). Seeing the Whole Elephant: Systematically Understanding and Uncovering Evaluation Biases in Automated Program Repair. *ACM TOSEM*, DOI: 10.1145/3561382.
20. **Yi, X., Chen, L., Mao, X., & Ji, T.** (2019). Efficient automated repair of high floating-point errors in numerical libraries. *ACM FSE*, DOI: 10.1145/3290369.
21. **Zhang, Q., Fang, C., Ma, Y., Sun, W., & Chen, Z.** (2023). A Survey of Learning-based Automated Program Repair. *ACM Computing Surveys*, DOI: 10.1145/3631974.
22. **Zhang, Q., Fang, C., Xie, Y., Ma, Y., Sun, W., Yang, Y., & Chen, Z.** (2026). A Systematic Literature Review on Large Language Models for Automated Program Repair. *ACM Computing Surveys*, DOI: 10.1145/3799693.
23. **Zhang, Y., Costea, A., Shariffdeen, R., McCall, D., & Roychoudhury, A.** (2025). EffFix: Efficient and Effective Repair of Pointer Manipulating Programs. *ACM TOPLAS*, DOI: 10.1145/3705310.
24. **Zhao, Y., Su, T., Liu, Y., Zheng, W., Wu, X., Kavuluru, R., Halfond, W. G. J., & Yu, T.** (2022). ReCDroid+: Automated End-to-End Crash Reproduction from Bug Reports for Android Apps. *ACM TOSEM*, DOI: 10.1145/3488244.
25. **Zirak, A., & Hemmati, H.** (2024). Improving Automated Program Repair with Domain Adaptation. *ACM TOSEM*, DOI: 10.1145/3631972.